

Mitigating the Apache Killer Vulnerability

David Scholberg
CS Master's Project

What is the Apache Killer?

- A vulnerability involving how Apache processes HTTP range requests
- A specially crafted HTTP request can cause a DoS attack through excessive CPU and RAM usage
- Versions 2.2.19 and prior are vulnerable

An example HTTP header

GET / HTTP/1.1

Host: www.google.com

Range: bytes=12-294,150-304

Accept-Encoding: gzip

Connection: keep-alive

HTTP header that may kill Apache

GET / HTTP/1.1

Host: www.google.com

Range:bytes=0-1,0-2,0-3,0-4,0-5,0-6,0-7,0-8,0-9,0-10,0-11,0-12,0-13,0-14,0-15,0-16,0-17,0-18,0-19,0-20,0-21,0-22,0-23,0-24,0-25,0-26,0-27,0-28,0-29,0-30,0-31,0-32,0-33,0-34,0-35,0-36,0-37,0-38,0-39,0-40,0-41,0-42,0-43,0-44,0-45,0-46,0-47,0-48,0-49,0-50,0-51,0-52,0-53,0-54,0-55,0-56,0-57,0-58,0-59,0-60,0-61,0-62,0-63,0-64,0-65,0-66,0-67,0-68,0-69,0-70,0-71,0-72,0-73,0-74,0-75,0-76,0-77,0-78,0-79,0-80,0-81,0-82,0-83,0-84,0-85,0-86,0-87,0-88,0-89,0-90,0-91,0-92,0-93,0-94,0-95,0-96,0-97,0-98,0-99,0-100,0-101,0-102

Accept-Encoding: gzip

Connection: keep-alive

The vulnerability

- Vulnerable versions of Apache process overlapping ranges very inefficiently
- An HTTP header with too many overlapping ranges will cause a DoS through RAM and CPU consumption
- Such an HTTP header is fairly trivial to construct and send

The fix

- Create an Intrusion Prevention System (IPS)
- IPS is a proxy server in front of Apache
- Traffic destined for Apache goes to the proxy server first
- Proxy server analyzes HTTP header; if too many overlapping ranges, block connection and log IP address

Caveat

- Apache Killer was fixed very quickly
- Has long since been fixed
- Apache Killer is used to illustrate the process of creating an IPS for application level exploits

IPS Design

IPS is a special Intrusion Detection System (IDS) that detects and blocks attacks

Must be in between client and server to block attacks

Can operate on any level of network stack, and usually multiple levels

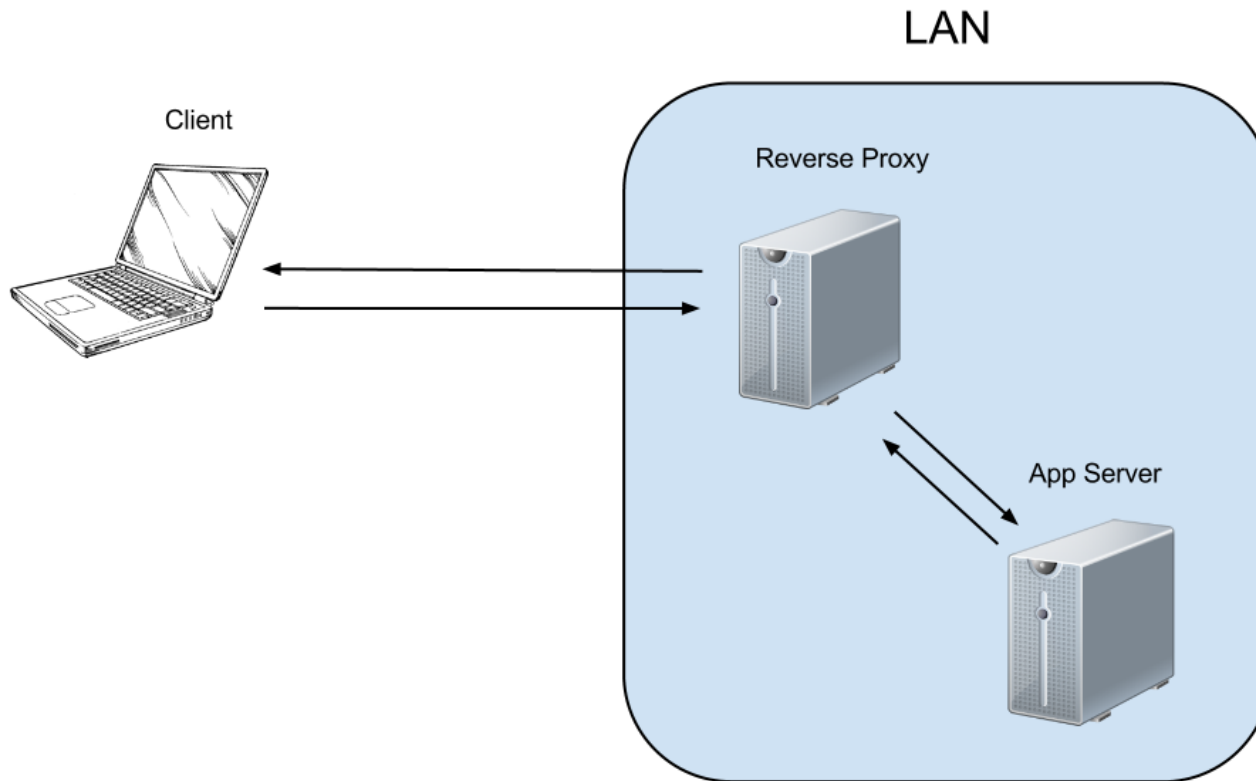
Reverse Proxy

IPS will be designed as a reverse proxy

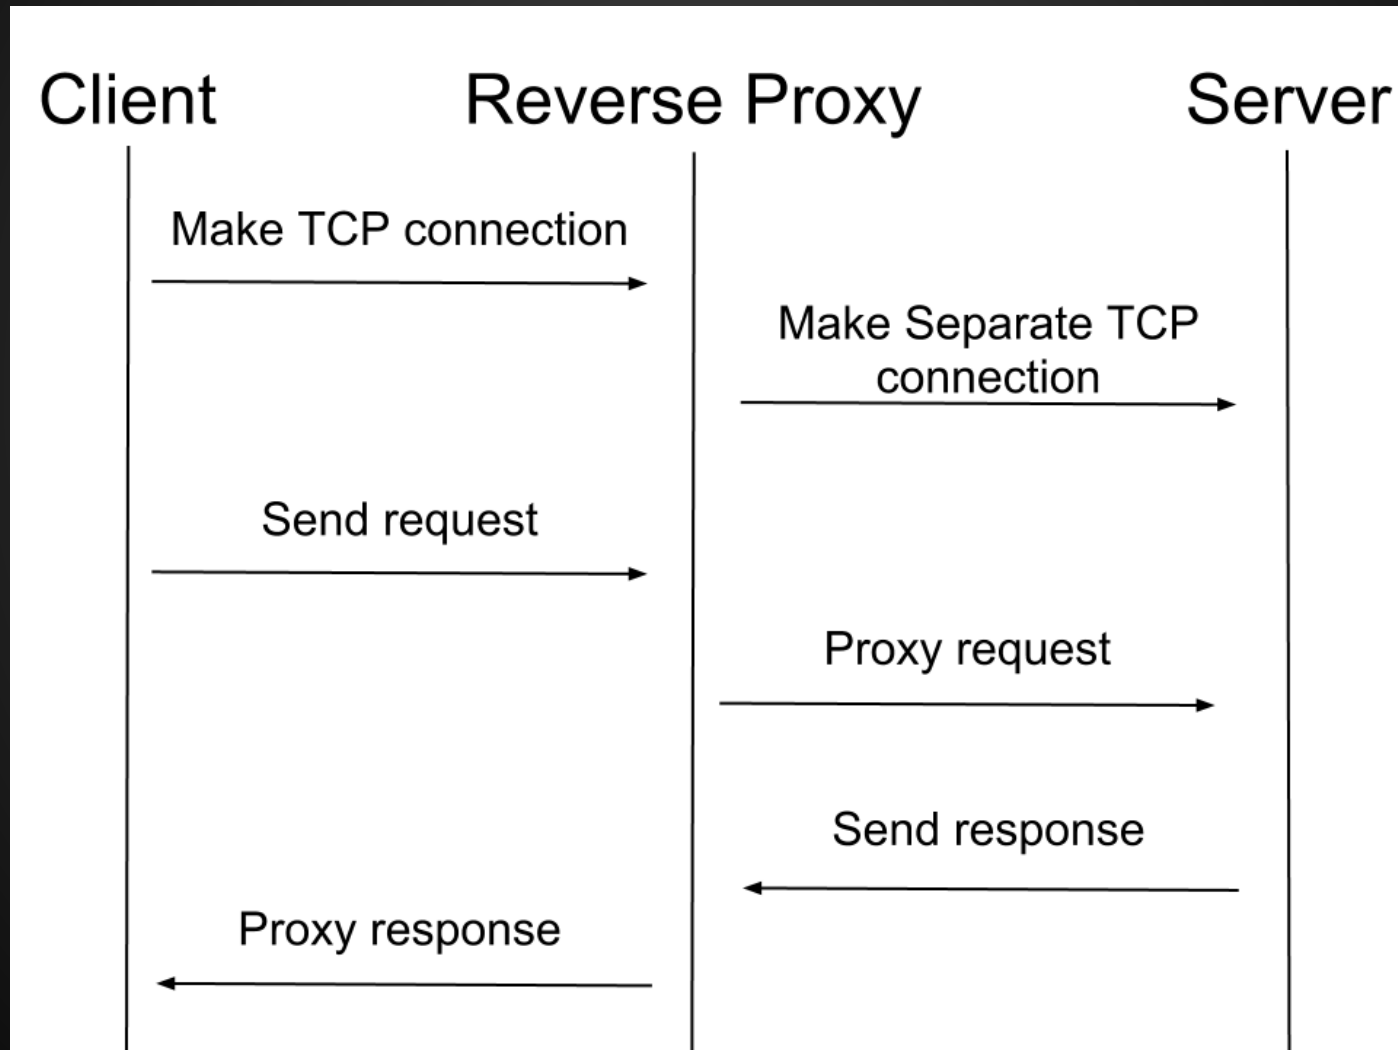
Reverse proxy is a server that proxies traffic unbeknownst to the client

Reverse Proxy

Example Reverse Proxy



Reverse Proxy



Reverse Proxy Pseudocode 1

```
while true
  read from client socket into free space in client buffer
  if more client data was received
    determine fate of client data (allow, buffer, or block)
    if client data is blocked
      exit loop

  read from server socket into free space in server buffer
  if more server data was received
    determine fate of server data (allow, buffer, or block)
    if server data is blocked
      exit loop
```

Reverse Proxy Pseudocode 2

```
if there is client data to send and data is not to be buffered
    send client data to server
    if not all data was sent
        shift remaining client data to start of client buffer

if there is server data to send and data is not to be buffered
    send server data to client
    if not all data was sent
        shift remaining server data to start of server buffer

end while
```

Apache Killer IPS

After client data is received, but before it is sent along to the server, it is examined

If data from the client contains too many ranges, drop data and close connections

Data from the server is implicitly trusted

Apache Killer IPS

Basic algorithm for attack detection:

```
set proxy flag to allow
```

```
if client data doesn't contain "\r\n\r\n"
```

```
    set proxy flag to buffer
```

```
else if client data contains valid HTTP declaration
```

```
    if client data has a range section
```

```
        if number of ranges exceeds threshold
```

```
            set proxy flag to block and log client IP
```

Implementation

Written in C for GNU/Linux

Linux socket library used for TCP connections

Perl Compatible Regular Expressions (PCRE)
library used for pattern matching

Overview of Linux socket programming

A listening socket can be created with calls to `socket()`, `bind()`, and `listen()`

Client connections can be accepted by calling `accept()` on the listening socket, which will return new socket connected to client

The new socket can be used for read/write operations to communicate with client

Overview of Linux socket programming

To initiate a connection with a server, call `socket()` and `connect()`

Forking

```
// Create child process to handle client connection.
// The return value of fork tells the process whether it is the
// parent or the child.

int pid = fork();
if(pid == 0) {
    // this is the child process
    handle_client_socket(
        client_socket,
        inet_ntoa(client_addr.sin_addr)
    );
}
else if (pid < 0) {
    // fork() returned an error
    die("fork() failed");
}
else {
    // this is the parent process
    close(client_socket);
}
```

Regular Expressions

```
const char *regex_strings[REGEX_NUM] = {  
    "\r\n\r\n$",  
    "HTTP\[0-9]+\?\.[0-9]+\?[0-9]",  
    "Range: .+?=(.*?)\r",  
    ",?[0-9]*?\-[0-9]*"  
};
```

Logging with syslog

Setup logging

```
openlog("reverse_proxy", LOG_PID, LOG_USER);
```

Write log entry

```
syslog(  
    LOG_WARNING,  
    "data from %s was rejected\n",  
    client_addr  
);
```

Testing Environment

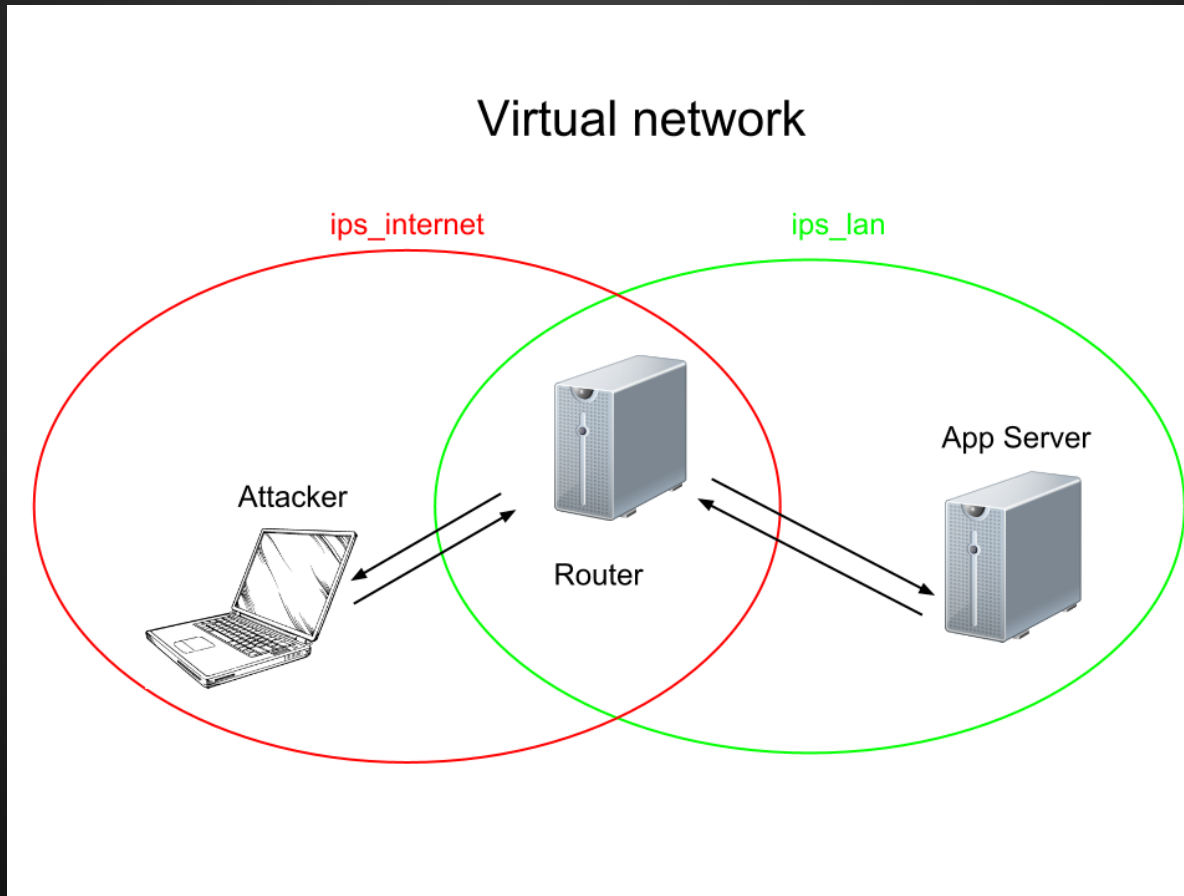
3 virtual machines (VMs), created with VirtualBox: Attacker, Router, and App Server

All 3 VMs run CentOS 6.3

The Router is on two virtual networks. The attacker is on one of those virtual networks, and the App Server is on the other.

Testing Environment

Network of VMs



Testing Environment

VM 1

Name :	Attacker
OS:	CentOS 6.3
RAM:	1024 MB
CPUs:	1
Internal Network:	ips_internet (IP: 7.0.0.2)
Software:	Metasploit, killapache.pl, w3m

Testing Environment

VM 2

Name :	Router
OS:	CentOS 6.3
RAM:	512 MB
CPUs:	2
Internal Network:	ips_internet (7.0.0.1), ips_lan (10.0.0.1)
Software:	iptables, Apache Killer IPS

Testing Environment

VM 3

Name:	App Server
OS:	CentOS 6.3
RAM:	3072 MB
CPUs:	2
Internal Network:	ips_lan (IP: 10.0.0.2)
Software:	Apache 2.2.18

Results - No IPS

Before Attack

```
top - 16:40:46 up 3 min, 1 user, load average: 0.01, 0.02, 0.00
Tasks: 86 total, 1 running, 85 sleeping, 0 stopped, 0 zombie
Cpu(s): 0.0%us, 0.2%sy, 0.0%ni, 99.8%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 2956912k total, 143492k used, 2813420k free, 5736k buffers
Swap: 1015800k total, 0k used, 1015800k free, 27260k cached
```

PID	USER	PR	NI	UIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
9	root	20	0	0	0	0	S	0.3	0.0	0:00.56	ksoftirqd/1
1010	root	20	0	15016	1264	984	R	0.3	0.0	0:00.34	top
1	root	20	0	19228	1496	1228	S	0.0	0.1	0:01.06	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.01	kthreadd
3	root	RT	0	0	0	0	S	0.0	0.0	0:00.05	migration/0
4	root	20	0	0	0	0	S	0.0	0.0	0:00.04	ksoftirqd/0
5	root	RT	0	0	0	0	S	0.0	0.0	0:00.00	migration/0
6	root	RT	0	0	0	0	S	0.0	0.0	0:00.00	watchdog/0
7	root	RT	0	0	0	0	S	0.0	0.0	0:01.01	migration/1
8	root	RT	0	0	0	0	S	0.0	0.0	0:00.00	migration/1
10	root	RT	0	0	0	0	S	0.0	0.0	0:00.00	watchdog/1
11	root	20	0	0	0	0	S	0.0	0.0	0:00.04	events/0
12	root	20	0	0	0	0	S	0.0	0.0	0:00.20	events/1
13	root	20	0	0	0	0	S	0.0	0.0	0:00.00	cgroup
14	root	20	0	0	0	0	S	0.0	0.0	0:00.00	khelper
15	root	20	0	0	0	0	S	0.0	0.0	0:00.00	netns
16	root	20	0	0	0	0	S	0.0	0.0	0:00.00	async/mgr
17	root	20	0	0	0	0	S	0.0	0.0	0:00.00	pm
18	root	20	0	0	0	0	S	0.0	0.0	0:00.00	sync_supers
19	root	20	0	0	0	0	S	0.0	0.0	0:00.00	bdi-default
20	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kintegrityd/0
21	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kintegrityd/1
22	root	20	0	0	0	0	S	0.0	0.0	0:00.15	kblockd/0
23	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kblockd/1
24	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kacpid
25	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kacpi_notify
26	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kacpi_hotplug
27	root	20	0	0	0	0	S	0.0	0.0	0:00.01	ata/0
28	root	20	0	0	0	0	S	0.0	0.0	0:00.00	ata/1
29	root	20	0	0	0	0	S	0.0	0.0	0:00.00	ata_aux

Results - No IPS

After Attack

```
top - 16:42:36 up 5 min, 1 user, load average: 16.62, 4.08, 1.35
Tasks: 126 total, 1 running, 125 sleeping, 0 stopped, 0 zombie
Cpu(s): 6.4%us, 14.5%sy, 0.0%ni, 0.0%id, 70.7%wa, 2.3%hi, 6.1%si, 0.0%st
Mem: 2956912k total, 2898520k used, 58392k free, 92k buffers
Swap: 1015800k total, 971996k used, 43804k free, 2804k cached
```

PID	USER	PR	NI	UIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
22	root	20	0	0	0	0	S	3.8	0.0	0:00.96	kblockd/0
1044	daemon	20	0	113m	86m	140	D	3.1	3.0	0:00.66	httpd
1042	daemon	20	0	118m	92m	140	D	2.8	3.2	0:00.61	httpd
9	root	20	0	0	0	0	S	2.5	0.0	0:01.07	ksoftirqd/1
1015	daemon	20	0	160m	55m	12	D	2.5	1.9	0:00.92	httpd
1024	daemon	20	0	149m	86m	136	D	2.2	3.0	0:00.75	httpd
1027	daemon	20	0	113m	76m	136	D	2.2	2.7	0:00.63	httpd
1041	daemon	20	0	118m	81m	140	D	2.2	2.8	0:00.62	httpd
12	root	20	0	0	0	0	S	2.0	0.0	0:00.45	events/1
1014	daemon	20	0	160m	56m	12	D	2.0	2.0	0:00.93	httpd
1019	daemon	20	0	140m	75m	136	D	2.0	2.6	0:00.75	httpd
1036	daemon	20	0	114m	77m	136	D	2.0	2.7	0:00.62	httpd
4	root	20	0	0	0	0	S	1.9	0.0	0:00.23	ksoftirqd/0
1018	daemon	20	0	131m	67m	136	D	1.9	2.3	0:00.70	httpd
1035	daemon	20	0	119m	81m	136	D	1.9	2.8	0:00.60	httpd
1039	daemon	20	0	109m	74m	140	D	1.9	2.6	0:00.59	httpd
1046	daemon	20	0	98.8m	71m	140	D	1.9	2.5	0:00.52	httpd
1047	daemon	20	0	114m	87m	140	D	1.9	3.0	0:00.62	httpd
1021	daemon	20	0	130m	66m	136	D	1.7	2.3	0:00.66	httpd
1023	daemon	20	0	127m	66m	136	D	1.7	2.3	0:00.67	httpd
1025	daemon	20	0	133m	71m	136	D	1.7	2.5	0:00.73	httpd
1029	daemon	20	0	115m	78m	136	D	1.7	2.7	0:00.60	httpd
1033	daemon	20	0	120m	83m	136	D	1.7	2.9	0:00.59	httpd
1040	daemon	20	0	113m	78m	140	D	1.7	2.7	0:00.58	httpd
1043	daemon	20	0	99964	71m	140	D	1.7	2.5	0:00.52	httpd
38	root	20	0	0	0	0	D	1.6	0.0	0:00.47	kswapd0
1016	daemon	20	0	160m	60m	12	D	1.6	2.1	0:00.86	httpd
1017	daemon	20	0	160m	53m	12	D	1.6	1.8	0:00.84	httpd
1028	daemon	20	0	106m	70m	136	D	1.6	2.5	0:00.58	httpd
1031	daemon	20	0	103m	68m	136	D	1.6	2.4	0:00.57	httpd

Results - With IPS

When exploit is run, there is no noticeable increase in memory consumption or CPU load on the App Server.

Legitimate client requests, made by w3m on the Attacker's machine, are allowed through and their responses are sent back without error

Results - With IPS

Excerpt of log output

```
Dec  6 02:39:21 ipsrouter reverse_proxy[1313]: data from 7.0.0.2 was rejected
Dec  6 02:39:21 ipsrouter reverse_proxy[1314]: data from 7.0.0.2 was buffered
Dec  6 02:39:21 ipsrouter reverse_proxy[1314]: data from 7.0.0.2 was rejected
Dec  6 02:39:21 ipsrouter reverse_proxy[1315]: data from 7.0.0.2 was rejected
Dec  6 02:39:21 ipsrouter reverse_proxy[1316]: data from 7.0.0.2 was rejected
Dec  6 02:39:21 ipsrouter reverse_proxy[1317]: data from 7.0.0.2 was rejected
Dec  6 02:39:21 ipsrouter reverse_proxy[1318]: data from 7.0.0.2 was rejected
Dec  6 02:39:21 ipsrouter reverse_proxy[1319]: data from 7.0.0.2 was rejected
Dec  6 02:39:21 ipsrouter reverse_proxy[1320]: data from 7.0.0.2 was rejected
Dec  6 02:39:21 ipsrouter reverse_proxy[1321]: data from 7.0.0.2 was rejected
```

Conclusion

It works!

Conclusion

Improvements

- send warning email to sys admin
- decrease verbosity of output
- log all traffic from IP address of attacker
- add HTTPS support

Conclusion

Caveats to custom IPS

- should only be temporary and last resort
 - use official patch or workaround if available
- Custom IPS may have its own vulnerabilities
 - better to use tried and tested solution
 - ModSecurity

Conclusion

This project has expanded my knowledge of

- C programming in a Linux environment
- network protocols (TCP/IP, HTTP)
- system and network administration
- protecting against application layer remote exploits

Questions?

???